

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy import stats
from scipy.stats import skew, kurtosis, mode #Python libraries inferential statistics
import seaborn as sns #This is for generating Histogram with Ker
```

```
df = pd.read_csv('hotel_books.csv') #read the 'hotel_books.csv' file
df.head(5)
```

```
↗
```

| | day | clients | total_bill |
|---|-----|---------|------------|
| 0 | 1 | 33 | 23958 |
| 1 | 2 | 25 | 26812 |
| 2 | 3 | 5 | 24871 |
| 3 | 4 | 17 | 17954 |
| 4 | 5 | 28 | 29416 |

```
df.dtypes #check for data types
```

```
↗
```

| day | int64 |
|------------|-------|
| clients | int64 |
| total_bill | int64 |

dtype: object

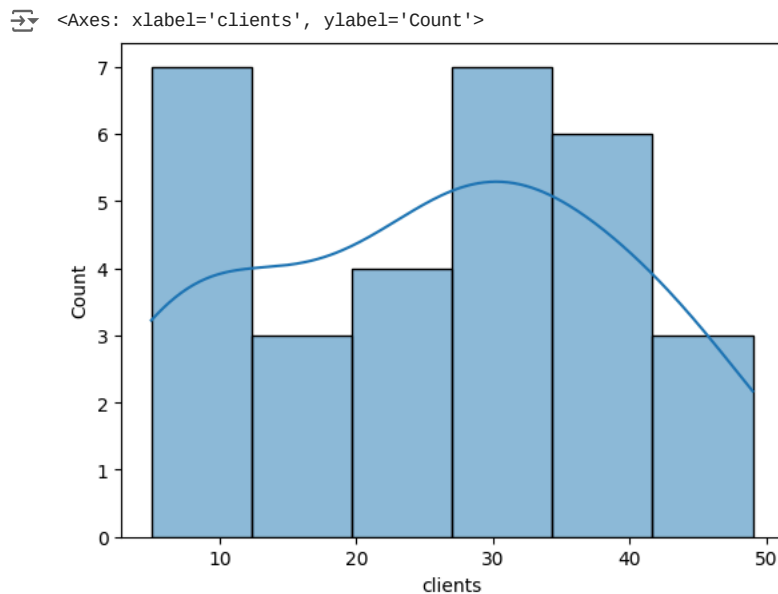
```
df.isnull().sum() #check for missing values
```

```
↗
```

| day | 0 |
|------------|---|
| clients | 0 |
| total_bill | 0 |


dtype: int64

```
sns.histplot(df['clients'], kde=True) #generate histogram with kernel density estimate (KDE) for number of hotel clients
```



```
sns.histplot(df['total_bill'], kde=True) #generate histogram with kernel density estimate (KDE) for total bill collected
```

<Axes: xlabel='total_bill', ylabel='Count'>



```
#compute for skewness and kurtosis for number of clients
skew1 = df['clients'].skew()
kurt1 = df['clients'].kurt()
print(f'Kurtosis for the number of hotel clients in a day:{kurt1}')
print(f'Skewness for the number of hotel clients in a day:{skew1}')
```

Kurtosis for the number of hotel clients in a day:-1.1388703400867874
Skewness for the number of hotel clients in a day:-0.05968808896371035

```
#compute for skewness and kurtosis for total number of bill
skew2 = df['total_bill'].skew()
kurt2 = df['total_bill'].kurt()
print(f'Kurtosis for the total bill collected in a day:{kurt2}')
print(f'Skewness for the total bill collected in a day:{skew2}')
```

Kurtosis for the total bill collected in a day:-1.130219880444574
Skewness for the total bill collected in a day:0.18976914965853053

df.describe() #generate summary measure and observe the mean and 50% (median)

| | day | clients | total_bill |
|-------|-----------|-----------|--------------|
| count | 30.000000 | 30.000000 | 30.000000 |
| mean | 15.500000 | 25.666667 | 28344.233333 |
| std | 8.803408 | 13.557879 | 12441.769892 |
| min | 1.000000 | 5.000000 | 7534.000000 |
| 25% | 8.250000 | 16.000000 | 18335.000000 |
| 50% | 15.500000 | 28.000000 | 25841.500000 |
| 75% | 22.750000 | 35.750000 | 39810.250000 |
| max | 30.000000 | 49.000000 | 49450.000000 |

```
stats.mode(df['clients']) #compute for mode
```

ModeResult(mode=8, count=4)

```
stats.mode(df['clients']) #compute for mode
```

ModeResult(mode=8, count=4)

